

# HW2

Yan Liu

9/15/2019

1.

a)

Let  $X$  be an exponential r.v. with rate 1. Using cdf inversion, write a function that generates  $n$  independent samples of  $X$  (we discussed this example in class, but you should do the cdf inversion yourself, not just quote our result). Compare the speed of your sampler for  $n = 10^6$  with that of your language's exponential sampler (in R `rexp`).

**Solution:**

Since  $X \sim \text{Exp}(1)$ , we have PDF  $f(x) = e^{-x}$ ,  $x \in [0, \infty)$ , and CDF  $F(x) = P(X \leq x) = \int_0^x e^{-t} dt = 1 - e^{-x}$ . Thus,  $x = F(F^{-1}(x)) = 1 - e^{-F^{-1}(x)}$ . Solve for  $F^{-1}(x)$ , we have  $F^{-1}(x) = -\ln(1 - x)$

```
time.start <- proc.time()
r.exp <- rexp(n=10^6)
proc.time() - time.start
```

```
##    user  system elapsed
##  0.043   0.002   0.045
```

```
myExpFunc <- function(N){
  return(-log(1-runif(N)))
}
time.start <- proc.time()
my.exp <- myExpFunc(N=10^6)
proc.time() - time.start
```

```
##    user  system elapsed
##  0.038   0.003   0.041
```

b)

Let  $X$  be a r.v. with the following pdf  $f(x) = \frac{1}{2} \exp(-|x|)$ , where  $\exp(-|x|) = e^{-|x|}$ . Write a function that samples from  $X$  using cdf inversion. Either theoretically or through numerical examples, show that your sampler correctly samples  $X$ .

**Solution**

$$f(x) = \frac{1}{2} \exp(-|x|) = \begin{cases} \frac{1}{2} e^{-x}, & \text{if } x \geq 0 \\ \frac{1}{2} e^x, & \text{otherwise} \end{cases}$$

Thus, when  $c < 0$ ,

$$F(x) = \Pr(X \leq c) = \int_{-\infty}^c \frac{1}{2} e^x dx = \frac{1}{2} e^c$$

When  $c \geq 0$ ,

$$\begin{aligned}
F(x) &= \Pr(x \leq c) \\
&= P(x < 0) + P(0 \leq x \leq c) \\
&= \int_{-\infty}^0 \frac{1}{2}e^x dx + \int_0^c \frac{1}{2}e^{-x} dx \\
&= \frac{1}{2}e^0 + (-\frac{1}{2}e^{-x}|_0^c) \\
&= 1 - \frac{1}{2}e^{-c}
\end{aligned}$$

Thus,

$$F(x) = \begin{cases} \frac{1}{2}e^x, & \text{if } x < 0 \\ 1 - \frac{1}{2}e^{-x}, & \text{otherwise} \end{cases}$$

Let  $y_1 = F(x)$  when  $x < 0$ ,  $y_2 = F(x)$  when  $x \geq 0$ .

$$\begin{aligned}
y_1 = \frac{1}{2}e^x &\implies x = \log(2y_1), \text{ where } y_1 < \frac{1}{2} \\
y_2 = 1 - \frac{1}{2}e^{-x} &\implies x = -\log(2 - 2y_2), \text{ where } y_2 \geq \frac{1}{2}
\end{aligned}$$

Thus,

$$F^{-1}(x) = \begin{cases} \log(2x), & \text{if } x < \frac{1}{2} \\ -\log(2 - 2x), & \text{otherwise} \end{cases}$$

Prove the two CDF are equivalent. Let  $C_x$  be any constant less than 0,  $y$  follows the distribution determined by  $F^{-1}$  above,  $C_y = \frac{1}{2}e^{C_x} < \frac{1}{2}$ , then,

$$\begin{aligned}
F_y(C_y) &= \Pr(y \leq C_y) \\
&= \Pr(F_y^{-1}(y) \leq F_y^{-1}(C_y)) \\
&= \Pr(x \leq \log(2 \times \frac{1}{2}e^{C_x})) \\
&= \Pr(x \leq C_x) \\
&= F_x(C_x)
\end{aligned}$$

Let  $C_x$  be any constant equal to or greater than 0,  $C_y = 1 - \frac{1}{2}e^{-C_x} \geq \frac{1}{2}$ , then,

$$\begin{aligned}
F_y(C_y) &= \Pr(y \leq C_y) \\
&= \Pr(F_y^{-1}(y) \leq F_y^{-1}(C_y)) \\
&= \Pr(x \leq -\log(2 - 2 \times (1 - \frac{1}{2}e^{-C_x}))) \\
&= \Pr(x \leq -\log(e^{-C_x})) \\
&= \Pr(x \leq C_x) \\
&= F_x(C_x)
\end{aligned}$$

## 2.

Write a function **MarkovChain(P, s<sub>0</sub>, s)** that simulates a Markov chain  $X(t)$  until the first time the chain is in state  $s$ , assuming  $X(0) = s_0$ . The function should return the path of the chain from  $t = 0$  to when it “hits” state  $s$ . You may use your language’s discrete sampler (in R **sample**) or write your own.

```

#' A helper function to determine the next state using runif()
#' @param row.cumsum cumsum probabilities of current row
#' @param states list of states
#' @returns next state
roll.dice <- function(row.cumsum, states){
  # roll a dice
  dice <- runif(1)
  n <- length(states)
  # could be the initial state
  if (dice >= 0 & dice < row.cumsum[1])
    return(states[1])
  else{
    for(i in seq(1,n)){
      # could be one of the middle states
      if (i <= n-1){
        if (dice >= row.cumsum[i] & dice < row.cumsum[i+1]){
          return(states[i+1])
        }
      }else{
        # otherwise the last state
        return(states[n])
      }
    }
  }
}

```

```

#' A function MarkovChain(P, s0, s) that simulates a Markov chain X(t)
#' until the first time the chain is in state s, assuming X(0) = s0.
#' @param P a transition probability matrix
#' @param s0 initial state
#' @param s target state or end state
#' @details P must satisfy: a square matrix; row sum must be 1;
#' each entry is less than or equal to 1;
#' the row names and the column names of P are both states.
#' The maximum steps to take is limited to 10000
#' @returns a list of states, i.e.,
#' the path of the chain from t = 0 to when it "hits" state s.
MarkovChain <- function(P,s0,s)
{
  if (is.null(s0))
    stop("Initial state must be non-null")

  if (is.null(s))
    stop("End state must be non-null")

  if (is.null(P)) {
    stop("Transition matrix must be non-null")
  }else{
    if (nrow(P) == ncol(P) & all(P <= 1) & all(P >= 0) & rowSums(P) == 1){
      path <- list()
      states <- row.names(P)
      s.current <- which(row.names(P)==s0)
      path[1] <- s0
      row.cumsum <- cumsum(P[s.current,])
    }
  }
}

```

```

s.next <- roll.dice(row.cumsum, states)
numSteps <- 1
maxSteps <- 10000
while (s.next != s & numSteps <= maxSteps ){
  row.num.current <- which(rownames(P)==s.next)
  row.cumsum <- cumsum(P[row.num.current,])
  path[numSteps+1] <- s.next
  s.next <- roll.dice(row.cumsum, states)
  numSteps <- numSteps + 1
}
path[numSteps+1] <- s.next
}
else{
  stop("Transition matrix must be correctly specified")
}
}
return (path)
}

```

```

states <- c("worse", "bad", "okay", "good", "better")
P <- matrix(c(0,.2,.3,.4,.1,
             .5,.5,.0,0,0,
             .05,.05,.05,.05,.8,
             .5,.5,.0,0,0,
             0,.2,.3,.4,.1), nrow = 5, ncol = 5, byrow = TRUE,
            dimnames = list(states,states))
set.seed(10)
MarkovChain(P,"worse","better")

```

```

## [[1]]
## [1] "worse"
##
## [[2]]
## [1] "good"
##
## [[3]]
## [1] "worse"
##
## [[4]]
## [1] "okay"
##
## [[5]]
## [1] "better"

```

```

set.seed(6)
MarkovChain(P,"good","better")

```

```

## [[1]]
## [1] "good"
##
## [[2]]
## [1] "bad"
##
## [[3]]
## [1] "bad"

```

```
##
## [[4]]
## [1] "worse"
##
## [[5]]
## [1] "okay"
##
## [[6]]
## [1] "better"
```

### 3.

Chutes and Ladders is a popular children's game. Use your function from problem 2 and a Monte Carlo approach to compute  $E[L]$  where  $L$  is the average length, in terms of the number of die rolls, of a Chutes and Ladders game. Use a CLT argument to determine roughly how many games you have to simulate to estimate  $E[L]$  to an accuracy of  $\pm 5$ .

```
chute.ladder <- read.csv("chutes_and_ladder_locations.csv")
game.mx <- matrix(data = rep(0,100*100),
                  nrow = 100,
                  ncol = 100,
                  dimnames = list(seq(1,100),seq(1,100)))

for (i in 1:94){
  for (j in 1:6){
    game.mx[i, i+j] <- 1/6
  }
}

for (i in 95:99){
  j = 1
  while (j <= 99-i){
    game.mx[i,i+j] <- 1/6
    j <- j+1
  }
  game.mx[i,100] <- 1-sum(game.mx[i,1:99])
}
```

Take a glance

First few entries

```
game.mx[1:5,1:5]
```

```
##      1      2      3      4      5
## 1 0 0.1666667 0.1666667 0.1666667 0.1666667
## 2 0 0.0000000 0.1666667 0.1666667 0.1666667
## 3 0 0.0000000 0.0000000 0.1666667 0.1666667
## 4 0 0.0000000 0.0000000 0.0000000 0.1666667
## 5 0 0.0000000 0.0000000 0.0000000 0.0000000
```

## Last few entries

```
game.mx[95:100,95:100]
```

```
##      95      96      97      98      99      100
## 95    0 0.1666667 0.1666667 0.1666667 0.1666667 0.3333333
## 96    0 0.0000000 0.1666667 0.1666667 0.1666667 0.5000000
## 97    0 0.0000000 0.0000000 0.1666667 0.1666667 0.6666667
## 98    0 0.0000000 0.0000000 0.0000000 0.1666667 0.8333333
## 99    0 0.0000000 0.0000000 0.0000000 0.0000000 1.0000000
## 100   0 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
```

## Update based on chute and ladder locations

```
entrances <- chute.ladder$start
exits <- chute.ladder$end
update <- function(game.mx, entrances, exits){
  len <- length(entrances)
  updated.mx <- game.mx
  for (i in seq(1:len)){
    updated.mx[entrances[i],] <- updated.mx[exits[i],]
  }
  return(updated.mx)
}
update.mx <- update(game.mx, entrances,exits)
```

## Take a glance

### First ladder

```
update.mx[c(1,38),35:40]
```

```
##      35 36 37 38      39      40
## 1     0 0 0 0 0.1666667 0.1666667
## 38    0 0 0 0 0.1666667 0.1666667
```

### First chute

```
update.mx[c(98,78),80:85]
```

```
##      80      81      82      83      84 85
## 98 0.1666667 0.1666667 0.1666667 0.1666667 0.1666667 0
## 78 0.1666667 0.1666667 0.1666667 0.1666667 0.1666667 0
```

```
playGame <- function(N){
  steps = numeric()
  for (i in 1:N){
    dice.0 <- runif(1, min = 0, max = 6)
    s0 <- case_when(
      dice.0 <= 1~38,
      dice.0 > 1 & dice.0 <= 2~2,
```

```


        dice.0 > 2 & dice.0 <= 3~3,
        dice.0 > 3 & dice.0 <= 4~14,
        dice.0 > 4 & dice.0 <= 5~5,
        dice.0 > 5 & dice.0 <= 6~6
    )

    path <- MarkovChain(update.mx, s0,100)
    steps[i] <- length(path)
  }
  return(steps)
}
set.seed(1024)
# number of times
Ns <- seq(50, 1000, 50)
sample.means <- numeric()
sample.sigmas <- numeric()
for (i in 1:length(Ns)){
  steps <- playGame(Ns[i])
  sample.means[i] <- mean(steps)
  sample.sigmas[i] <- sqrt(var(steps))
}

```


Check convergence of sample mean

```
plot(Ns, sample.means, type = "b", xlab = "Number of trials")
```



Check convergence of sample standard deviation

```
plot(Ns,sample.sigmas, type = "b", xlab = "Number of trials")
```



Next, determine the number of trials that would enable the estimated  $\hat{E}[L]$  to have an accuracy of  $\pm 5$  compared with the true  $E[L]$ . Based on the central limit theorem,

$$\lim_{n \rightarrow \infty} \sqrt{n} \left( \frac{1}{n} \sum_{i=1}^n \hat{L}_i - E(L) \right) \sim N(0, \sigma^2)$$

Let

$$\hat{E}[L] = \frac{1}{n} \sum_{i=1}^n \hat{L}_i$$

In other words,

$$\Pr(-5 \leq \hat{E}[L] - E[L] \leq 5) = \Pr(\hat{E}[L] - 5 \leq E[L] \leq \hat{E}[L] + 5)$$

After some transformations, we will have:

$$\Pr\left(\frac{-5\sqrt{n}}{\sigma} \leq N(0, 1) \leq \frac{5\sqrt{n}}{\sigma}\right)$$

Assuming the confidence level is 99%, which corresponds to a Z-score of 2.58. Let  $\frac{5\sqrt{n}}{\sigma} = 2.58$ , where  $\sigma$  is estimated by  $\hat{\sigma}$ . According to the second plot above,  $\hat{\sigma} \approx 24$ . Solve the equation for  $n$ , we have  $n \approx 153$ .