

Homework 4 Yan Liu

Problem 1

(c)

```
library(dplyr, quietly = T)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##     filter, lag
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

library(magrittr)
# The function below is adapted from the HW3 solution
exp_logL <- function(X, theta)
{
  p1 <- theta["p1"]; p2 <- 1-p1
  mu1 <- theta["mu1"]; mu2 <- theta["mu2"]
  s1 <- theta["sigma1"]; s2 <- theta["sigma2"]

  N1 <- dnorm(X, mean=mu1, sd=sqrt(s1))
  N2 <- dnorm(X, mean=mu2, sd=sqrt(s2))
  D <- p1*N1 + p2*N2

  r1 <- p1*N1/D
  r2 <- 1-r1

  exp_ll <- (r1*log(p1*N1) + r2*log(p2*N2)) %>% sum
  return (exp_ll)
}

# test
theta <- c(mu1=70, mu2=60, sigma1=10, sigma2=15, p1=.3)
X <- read.table("Hope Heights.txt", header=T) %>% dplyr::pull(Height)

print(exp_logL(X, theta))

## [1] -378.3055

grad_logL <- function(X, theta)
{
  p1 <- theta["p1"]; p2 <- 1-p1
  mu1 <- theta["mu1"]; mu2 <- theta["mu2"]
  s1 <- theta["sigma1"]; s2 <- theta["sigma2"]

  N1 <- dnorm(X, mean=mu1, sd=sqrt(s1))
```

```

N2 <- dnorm(X, mean=mu2, sd=sqrt(s2))

D <- p1*N1 + p2*N2
r1 <- p1*N1/D
r2 <- 1-r1
r1.sum <- r1 %>% sum
r2.sum <- r2 %>% sum

N <- length(X)
grad_samples <- matrix(c(mu1=r1*X/r1.sum,
                          mu2=r2*X/r2.sum,
                          s1 = r1*(X-mu1)^2/r1.sum,
                          s2 = r2*(X-mu2)^2/r2.sum,
                          p1= r1.sum/N),
                       nrow=N,
                       ncol=length(theta))

grad <- grad_samples %>% colSums %>% setNames(c("mu1", "mu2", "s1", "s2", "p1"))
return (grad_samples %>% colSums)
}

# test
print(grad_logL(X, theta))

## Warning in matrix(c(mu1 = r1 * X/r1.sum, mu2 = r2 * X/r2.sum, s1 = r1 * :
## data length [401] is not a sub-multiple or multiple of the number of rows
## [100]
## [1] 71.01751 64.84814 10.88781 31.57688 70.80153

norm <- function(x) sqrt(sum(x^2))

expectation_maximization <- function(X, start_theta,
                                     max_iter=100)
{
  theta <- start_theta
  iter <- 1

  g <- grad_logL(X, theta)
  while (norm(g) > 1E-3 & iter <= max_iter) {

    if (iter %% 10 == 0)
      cat("iter=", iter, "likelihood=", exp_logL(X, theta), "\n")

    ng <- g/norm(g)
    s <- 1
    new_theta <- theta + s*ng

    # back up if we exceed boundaries
    while(new_theta["p1"] < 0 | new_theta["p1"] > 1 |
          new_theta["sigma1"] < 0 | new_theta["sigma2"] < 0) {
      s <- s/2
      new_theta <- theta + s*ng
    }
  }
}

```

```

# backtrack
while(exp_logL(X, new_theta) < exp_logL(X, theta)) {
  s <- s/2
  new_theta <- theta + s*ng
}

theta <- new_theta
g <- grad_logL(X, theta)
iter <- iter + 1
}

return (list(theta=theta,
             likelihood=exp_logL(X, theta)))
}

# try a few starting points
start_theta <- c(mu1=75, mu2=65, sigma1=5, sigma2=5, p1=.3)
out1 <- expectation_maximization(X, start_theta, max_iter=100)

```

```

## iter= 10 likelihood= -327.9474
## iter= 20 likelihood= -327.9474
## iter= 30 likelihood= -327.9474
## iter= 40 likelihood= -327.9474
## iter= 50 likelihood= -327.9474
## iter= 60 likelihood= -327.9474
## iter= 70 likelihood= -327.9474
## iter= 80 likelihood= -327.9474
## iter= 90 likelihood= -327.9474
## iter= 100 likelihood= -327.9474

```

```
print(out1)
```

```

## $theta
##      mu1      mu2      sigma1      sigma2      p1
## 75.1488221 65.1353336 5.0138827 5.0206773 0.4463439
##
## $likelihood
## [1] -327.9474

```

```

start_theta <- c(mu1=10, mu2=20, sigma1=20, sigma2=20, p1=.5)
out2 <- expectation_maximization(X, start_theta, max_iter=100)

```

```

## iter= 10 likelihood= -5076.949
## iter= 20 likelihood= -4208.267
## iter= 30 likelihood= -3823.811
## iter= 40 likelihood= -3823.811
## iter= 50 likelihood= -3823.811
## iter= 60 likelihood= -3823.811
## iter= 70 likelihood= -3823.811
## iter= 80 likelihood= -3823.811
## iter= 90 likelihood= -3823.811
## iter= 100 likelihood= -3823.811

```

```
print(out2)
```

```
## $theta
##      mu1      mu2      sigma1      sigma2      p1
## 10.4694963 20.5010002 41.8108540 37.5706493 0.9678372
##
## $likelihood
## [1] -3823.811

start_theta <- c(mu1=80, mu2=60, sigma1=5, sigma2=5, p1=.2)
out3 <- expectation_maximization(X, start_theta, max_iter=100)
```

```
## iter= 10 likelihood= -721.0258
## iter= 20 likelihood= -721.0258
## iter= 30 likelihood= -721.0258
## iter= 40 likelihood= -721.0258
## iter= 50 likelihood= -721.0258
## iter= 60 likelihood= -721.0258
## iter= 70 likelihood= -721.0258
## iter= 80 likelihood= -721.0258
## iter= 90 likelihood= -721.0258
## iter= 100 likelihood= -721.0258
```

```
print(out3)
```

```
## $theta
##      mu1      mu2      sigma1      sigma2      p1
## 80.5202222 60.4723872 5.3427470 5.3570546 0.7111318
##
## $likelihood
## [1] -721.0258
```

The first starting point gives the highest likelihood and seems to split the heights roughly as we would expect (although only because we know the answer).

(d)

Let's choose the mixture based on which of $P(z_i = 1 \mid \hat{X}_i, \theta)$ and $P(z_i = 2 \mid \hat{X}_i, \theta)$ is bigger.

```
# Get the theta that was best from b.i
theta <- out1$theta
theta
```

```
##      mu1      mu2      sigma1      sigma2      p1
## 75.1488221 65.1353336 5.0138827 5.0206773 0.4463439
```

```
Gender <- read.table("Hope Heights.txt", header=T) %>% dplyr::pull(Gender)
```

```
classify_gender <- function(X, theta)
{
  p1 <- theta["p1"]; p2 <- 1-p1
  mu1 <- theta["mu1"]; mu2 <- theta["mu2"]
  s1 <- theta["sigma1"]; s2 <- theta["sigma2"]

  N1 <- dnorm(X, mean=mu1, sd=sqrt(s1))
  N2 <- dnorm(X, mean=mu2, sd=sqrt(s2))

  D <- p1*N1 + p2*N2
```

```

r1 <- p1*N1/D
r2 <- 1-r1

# The first mixture has higher mu, so associate that with male gender
guess <- ifelse(r1*log(p1*N1) > r2*log(p2*N2), 1, 2)

return (guess)
}

guess_Gender <- classify_gender(X, theta)
correct <- sum(Gender==guess_Gender)/length(Gender)
print(correct)

```

```
## [1] 0.85
```

We get roughly 85% right, which is better than in Homework 3. Medium heights is harder to classify or guess than extreme heights.

```
data.frame(height=X, guess=guess_Gender, true=Gender) %>% dplyr::arrange(height) -> result
result %>% head(5)
```

```
##   height guess true
## 1     60     1    1
## 2     60     1    1
## 3     61     1    1
## 4     61     1    1
## 5     61     1    1
```

```
result[45:55,]
```


```
##   height guess true
## 45     69     1    1
## 46     69     1    1
## 47     69     1    1
## 48     69     1    2
## 49     69     1    2
## 50     69     1    2
## 51     69     1    2
## 52     70     1    1
## 53     70     1    1
## 54     70     1    1
## 55     70     1    2
```

```
result %>% tail(5)
```

```
##   height guess true
## 96     75     2    2
## 97     76     2    2
## 98     76     2    2
## 99     77     2    2
## 100    79     2    2
```

Problem 2

```
X_b <- read.csv("noisy_bits.csv")
image(t(X_b))
```



```
exp_b_logL <- function(X_b, theta_b)
{
  p1 <- theta_b["p1",]; p2 <- 1-p1
  mu1 <- theta_b["mu1",]; mu2 <- theta_b["mu2",]

  # mu1=.2;
  # mu2=.2; p1=.45;
  # p2 <- 1-p1

  R1 <- p1*apply(mu1^X_b*(1-mu1)^(1-X_b), 1, prod)
  R2 <- p2*apply(mu2^X_b*(1-mu2)^(1-X_b), 1, prod)
  r1 <- R1 / (R1+R2)
  r2 <- 1-r1

  exp_b_ll <- (r1*log(R1) + r2*log(R2)) %>% sum
  return (exp_b_ll)
}

# test
theta_b<- matrix(rep(c(.4, .2, p1=.45), each =10), nrow =3, byrow = T) %>%
  set_rownames(c("mu1", "mu2", "p1"))
theta_b

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
## mu1 0.40 0.40 0.40 0.40 0.40 0.40 0.40 0.40 0.40 0.40
## mu2 0.20 0.20 0.20 0.20 0.20 0.20 0.20 0.20 0.20 0.20
## p1  0.45 0.45 0.45 0.45 0.45 0.45 0.45 0.45 0.45 0.45
```

```

#theta_b
print(exp_b_logL(X_b, theta_b))

## [1] -4036.815

grad_b_logL <- function(X_b, theta_b)
{
  p1 <- theta_b["p1",]; p2 <- 1-p1
  mu1 <- theta_b["mu1",]; mu2 <- theta_b["mu2",]

  R1 <- p1*apply(mu1^X_b*(1-mu1)^(1-X_b), 1, prod)
  R2 <- p2*apply(mu2^X_b*(1-mu2)^(1-X_b), 1, prod)
  r1 <- as.vector(R1 / (R1+R2))
  r2 <- 1-r1

  r1.sum <- r1 %>% sum
  r2.sum <- r2 %>% sum

  N <- dim(X_b)[1]
  dimension <- dim(X_b)[2]
  grad_samples <- matrix(c(mu1=(t(r1) %*% as.matrix(X_b)),
                           mu2=(t(r2) %*% as.matrix(X_b)),
                           p1= rep(r1.sum/N, dimension)),
                        byrow = T,
                        ncol=dimension) %>% set_rownames(c("mu1", "mu2", "p1"))

  return (grad_samples)
}

# test
print(grad_b_logL(X_b, theta_b))

##           [,1]      [,2]      [,3]      [,4]      [,5]
## mu1 221.1760721 215.2652768 211.0252129 198.2710868 218.7350830
## mu2  38.8239279  37.7347232  39.9747871  35.7289132  36.2649170
## p1   0.8040731  0.8040731  0.8040731  0.8040731  0.8040731
##           [,6]      [,7]      [,8]      [,9]     [,10]
## mu1 228.2423173 211.1050655 214.3007194 205.4398974 209.2377088
## mu2  43.7576827  36.8949345  38.6992806  33.5601026  33.7622912
## p1   0.8040731  0.8040731  0.8040731  0.8040731  0.8040731

expectation_maximization_b <- function(X_b, start_theta_b,
                                       max_iter=200)
{
  theta_b <- start_theta_b
  iter <- 1
  s <- 1
  gradient <- grad_b_logL(X_b, theta_b)

  while (iter <= max_iter) {

    # if (iter %% 10 == 0)
    #   cat("iter=", iter, "likelihood=", exp_b_logL(X_b, theta_b), "\n")

    ng <- gradient/norm(gradient)

```

```

s <- 1
new_theta_b <- theta_b + s*ng

# back up if we exceed boundaries
while(any(new_theta_b["p1",]) < 0) {
  s <- s/2
  new_theta_b <- theta_b + s*ng
}

# backtrack

while((exp_b_logL(X_b, new_theta_b) < exp_b_logL(X_b, theta_b))) {
  s <- s/2
  new_theta_b <- theta_b + s*ng
}

theta_b <- new_theta_b
gradient <- grad_b_logL(X_b, theta_b)
iter <- iter + 1

return (list(theta=theta_b,
             likelihood=exp_b_logL(X_b, theta_b)))
}
}

#start_theta <- c(mu1=.4, mu2=.3, p1=.4)
start_theta_b<- matrix(rep(c(.4, .2, p1=.2), each =10), nrow =3, byrow = T) %>%
  set_rownames(c("mu1", "mu2", "p1"))
final_out <- expectation_maximization_b(X_b, start_theta_b, max_iter=200)

# Get the theta that was best from b.i
theta <- final_out$theta

classify_X <- function(X, theta)
{
  p1 <- theta["p1",];
  p2 <- 1-p1;
  mu1 <- theta["mu1",];
  mu2 <- theta["mu2",]

  R1 <- p1*apply(mu1^X_b*(1-mu1)^(1-X_b), 1, prod)
  R2 <- p2*apply(mu2^X_b*(1-mu2)^(1-X_b), 1, prod)
  r1 <- R1 / (R1+R2)
  r2 <- 1-r1

  guess <- ifelse((r1*log(R1) > r2*log(R2)), 0, 1)

  return (guess)
}

guess_X <- classify_X(X_b, theta)
X_b_label <- ifelse(apply(X_b, 1, sum)>5,1,0)

```



```
print(sum(guess_X == X_b_label)/500)
```

```
## [1] 0.87
```

The result is fairly good.